

Reinforcement Learning for Language Model Training

Polina Tsvilodub

Reinforcement Learning

RL4
LMT

LMS: Recap

Transformers, Training & Inference

- ▶ $LM : X \mapsto \Delta(S)$
- ▶ transformers use self-attention to offer and retrieve relevant information
 - stacked transformer blocks and multi-head attention increase capacity
- ▶ LMs are trained to predict the next word using cross-entropy loss (via teacher-forcing)
- ▶ decoding schemes are used for inference given a trained LM
 - different stochastic sampling regimes
- ▶ SOTA models exhibit ‘in-context learning’
- ▶ advanced prompting techniques might improve LLMs’ generalization performance



Making LLMs useful

Adaptation

- ▶ training a task-specific head on top of a model
 - e.g., span prediction layer on top of BERT with frozen BERT
 - on a dataset of ground truth input-output pairs for a particular task
- ▶ fine-tuning the model
 - further training part or entire model for a shorter time
 - on a dataset of ground truth input-output pairs for a particular task
- ▶ practical problem
 - training with standard supervision is impractical (data collection)
 - and inefficient (restricting “ground truth” to finite set of answers for open-ended tasks)
- ▶ **RL is the solution:** learn to achieve goal based on feedback from environment rather than direct demonstration of correct behaviour



Reinforcement learning


Flavors of machine learning

- ▶ Unsupervised learning
 - e.g., clustering
- ▶ discover patterns in unlabeled data
 - ‘given my inductive bias, what is the likely structure of the data?’
- ▶ Supervised learning
 - also self-supervised learning
 - aka behavioural cloning
- ▶ learn to output Y , given X , from labeled data
 - ‘do as I show you’
- ▶ learning from demonstration
- ▶ Reinforcement learning
 - trial-and-error learning
- ▶ learning from interaction / experience
 - ‘how do I optimally behave in order to maximize reward?’
 - or, ‘how do I optimally achieve my goal?’
 - most natural way of learning?
 - tightly connected to the way organisms behave (“pleasure maximizers”)


Reinforcement Learning: Overview

Introduction

- ▶ **Reinforcement Learning:** Computational formalisation of goal-directed learning and decision making under uncertainty
- ▶ **Goal:** Maximize rewards (by learning optimal behavior)
- ▶ **Basic building blocks:**
 - Agent
 - States
 - Actions
 - Transition function P
 - Reward
 - Policy




Associative RL



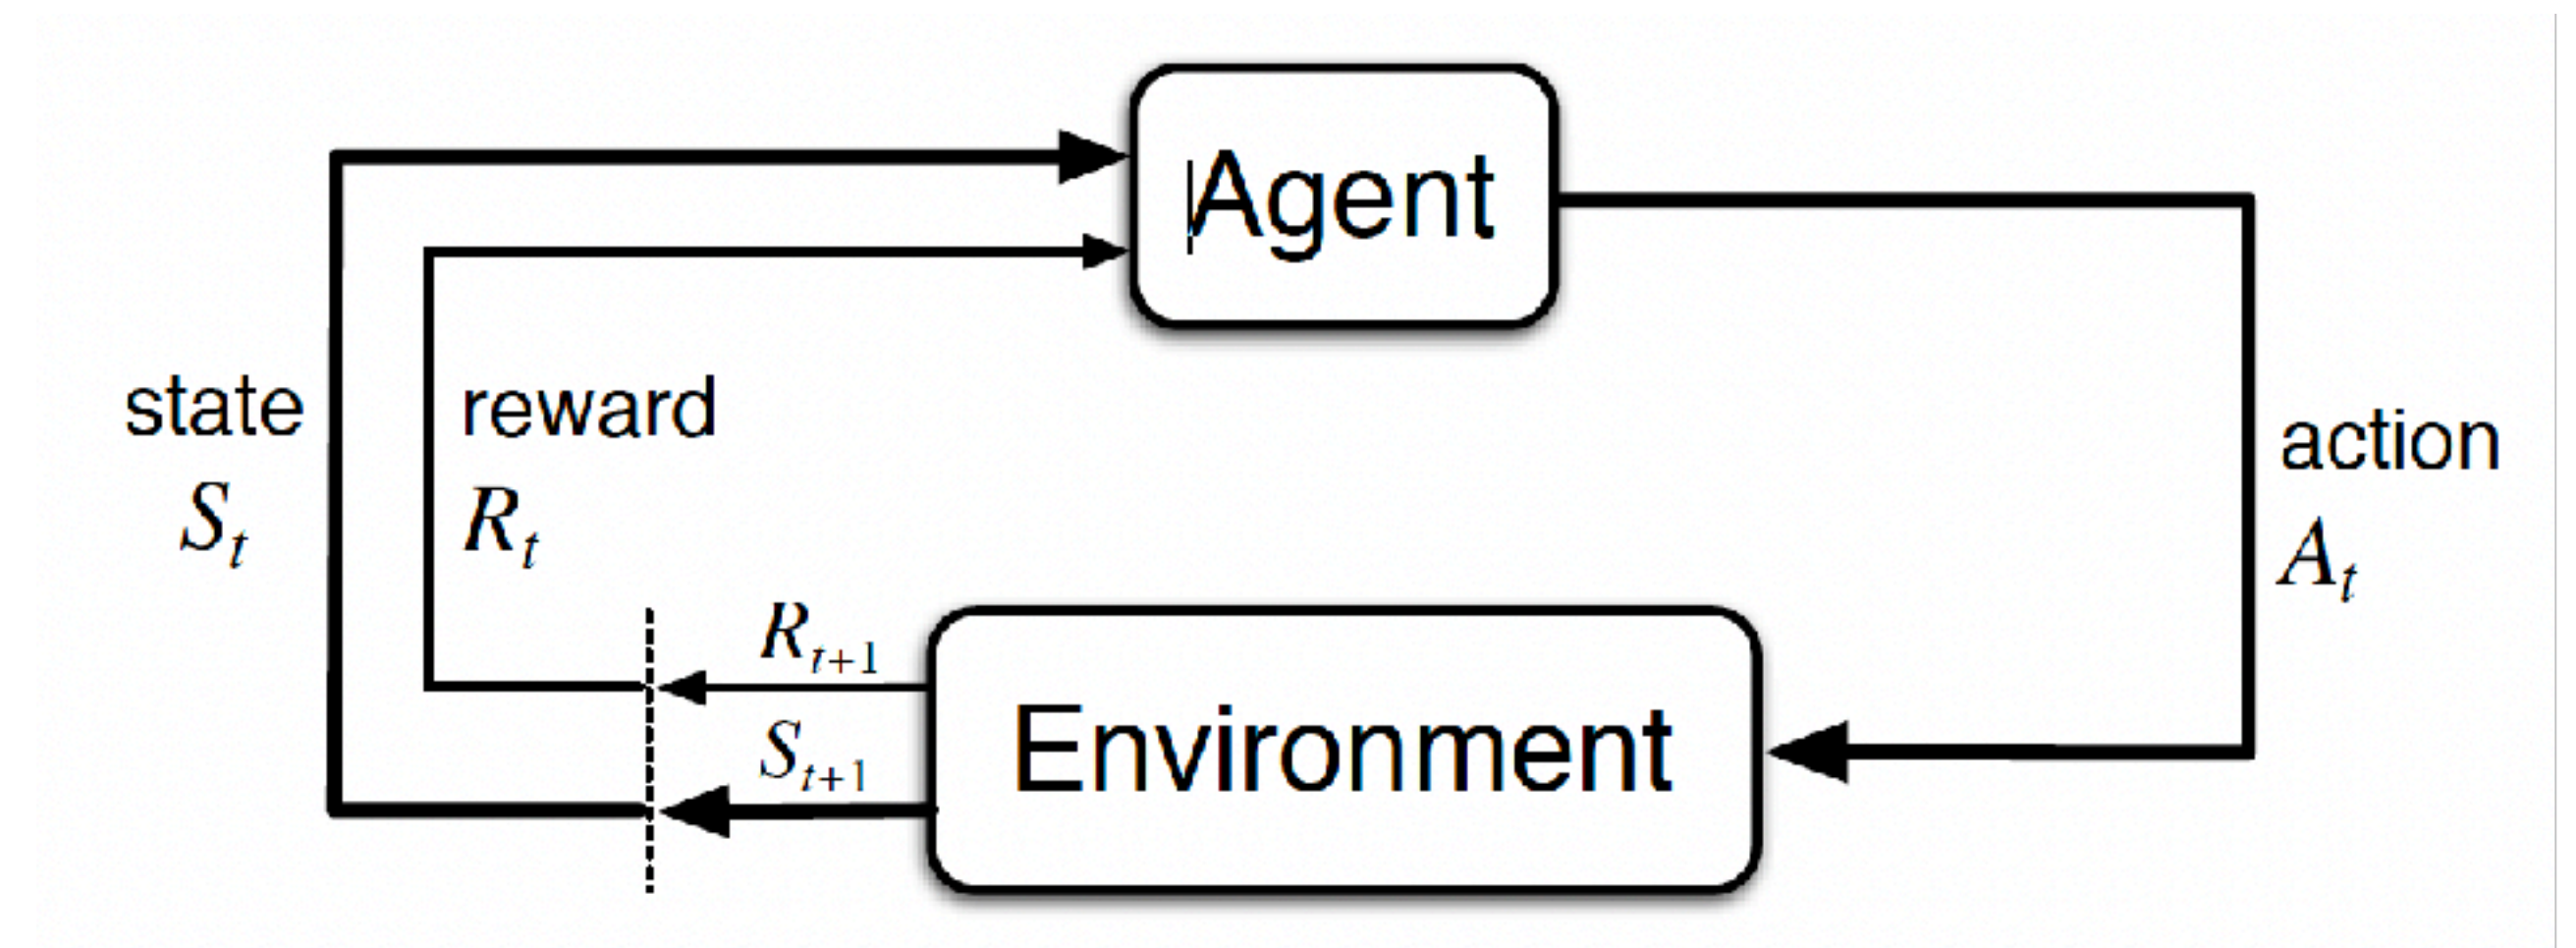
Reinforcement Learning: Overview

Introduction

- ▶ **Reinforcement Learning:** Computational formalisation of goal-directed learning and decision making under uncertainty
- ▶ **Goal:** Maximize rewards (by learning optimal behavior)
- ▶ **Basic building blocks:**
 - Agent
 - States: $S_t \in \mathcal{S}$ for $t = 0, 1, 2, 3, \dots$
 - Actions: $A_t \in \mathcal{A}(s)$
 - Transition function: $P(s' | s, a)$
 - Reward: $R_{t+1} \in \mathcal{R}$
 - Policy: $\pi(S_t) = P(A_t | S_t)$



Associative RL



Markov Decision Processes

Formal definition

Finite MDPs: (S, A, T, R)

1. $S_t \in S$ for $t = 0, 1, 2, 3, \dots$

2. $A_t \in A(s)$

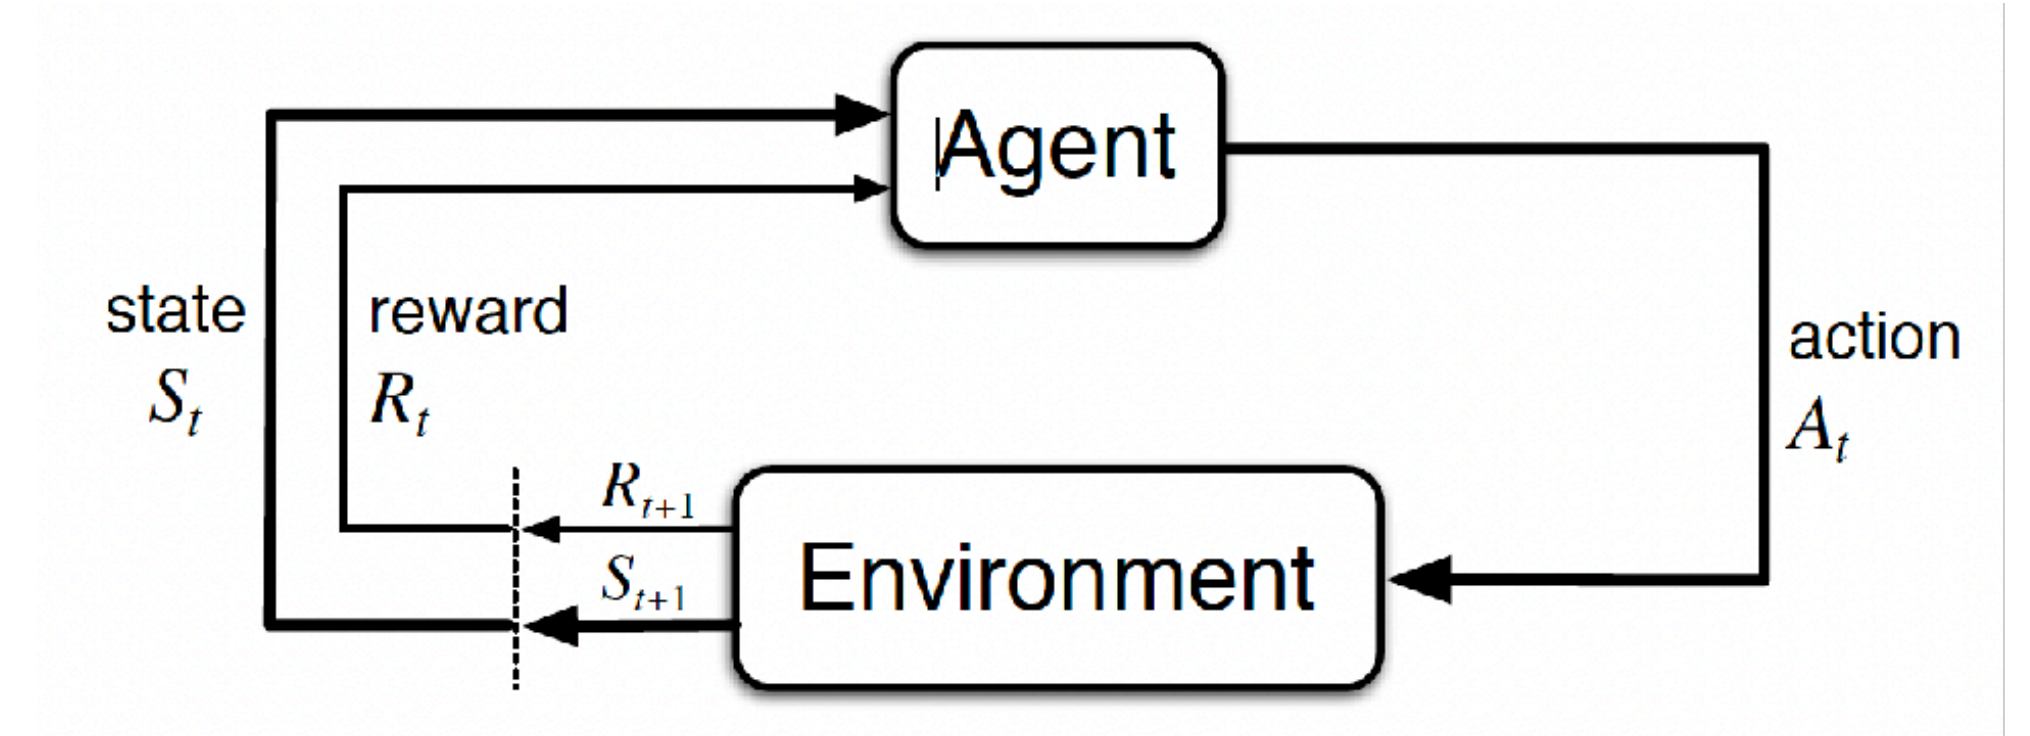
3. $R_{t+1} \in R$

4. $T(s' | s, a) = \sum_{r' \in R} P(s', r' | s, a)$

5. Horizon H , discount factor $0 \leq \gamma \leq 1$

Expected rewards for state s and action a : $r(s, a) = \mathbb{E}(R_t | S_{t-1} = s, A_{t-1} = a) = \sum_{r \in R} r \sum_{s' \in S} P(s', r | s, a)$

Markov property: S_{t-1}, A_{t-1} include all information about past agent-environment interactions that are relevant for S_t, R_t



Markov Decision Processes

Formal definition

Finite MDPs: (S, A, T, R)


1. $S_t \in S$ for $t = 0, 1, 2, 3, \dots$
2. $A_t \in A(s)$
3. $R_{t+1} \in R$
4. $T(s' | s, a) = \sum_{r' \in R} P(s', r' | s, a)$

Goal: maximize **returns** until goal achieved

$$G_t = R_{t+1} + R_{t_2} + \dots + R_T$$

Formally: maximize expected **discounted** rewards over **episode**

$$G_t = R_{t+1} + \gamma R_{t_2} + \gamma^2 R_{t+3} + \dots + \gamma^{T-t-1} R_T = \sum_{k=t+1}^T \gamma^{k-t-1} R_k$$



Markov Decision Processes


Formal definition

Finite MDPs: (S, A, T, R)

1. $S_t \in S$ for $t = 0, 1, 2, 3, \dots$
2. $A_t \in A(s)$
3. $R_{t+1} \in R$
4. $T(s' | s, a) = \sum_{r' \in R} P(s', r' | s, a)$

Goal: maximize discounted **returns**

$$G_t = R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \dots + \gamma^{T-t-1} R_T = \sum_{k=t+1}^T \gamma^{k-t-1} R_k$$
$$= R_{t+1} + \gamma G_{t+1}$$



- ▶ We can identify optimal way to behave if we know what good particular states and/or actions are:

Optimal **state-value function**:

$$V_{\pi}^*(s) = \max_{\pi} \mathbb{E}[G_t | S_t = s] = \max_{\pi} \mathbb{E}[R_{t+1} + \gamma G_{t+1} | S_t = s]$$
$$= \max_a \sum_{s', r}^{\pi} P(s', r | s, a) [r + \gamma G_{t+1} | S_t = s] \text{ for all } s$$

- ▶ Optimization problem: (computationally) find **optimal policy** $\pi^*(S_t) = P(A_t | S_t)$

Markov Decision Processes

Formal definition

Finite MDPs: (S, A, T, R)


1. $S_t \in S$ for $t = 0, 1, 2, 3, \dots$
2. $A_t \in A(s)$
3. $R_{t+1} \in R$
4. $T(s' | s, a) = \sum_{r' \in R} P(s', r' | s, a)$

Optimal state-value function:

$$V_{\pi}^*(s) = \max_{\pi} \mathbb{E}[G_t | S_t = s] = \max_{\pi} \mathbb{E}\left[\sum_{k=t+1}^T \gamma^{k-t-1} R_k | S_t = s\right]$$

deterministic optimal policy
 $\gamma = 1$

deterministic optimal policy
 $\gamma = 0.9$



$$\begin{aligned} V^*(4,3) &= 1 \\ V^*(3,3) &= 1 \\ V^*(2,3) &= 1 \\ V^*(1,1) &= 1 \\ V^*(4,2) &= -1 \end{aligned}$$

$$\begin{aligned} V^*(4,3) &= 1 \\ V^*(3,3) &= 0.9 \\ V^*(2,3) &= 0.81 \\ V^*(1,1) &= 0.9^5 = 0.59 \\ V^*(4,2) &= -1 \end{aligned}$$

Markov Decision Processes


Formal definition

Finite MDPs: (S, A, T, R)

1. $S_t \in S$ for $t = 0, 1, 2, 3, \dots$
2. $A_t \in A(s)$
3. $R_{t+1} \in R$
4. $T(s' | s, a) = \sum_{r' \in R} P(s', r' | s, a)$

Goal: maximize discounted **returns**

$$G_t = R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \dots + \gamma^{T-t-1} R_T = \sum_{k=t+1}^T \gamma^{k-t-1} R_k$$
$$= R_{t+1} + \gamma G_{t+1}$$



- We can identify optimal way to behave if we know what good particular states and/or actions are:

Optimal **action-value function**:

$$Q_{\pi}^*(s, a) = \max_{\pi} \mathbb{E}[G_t | S_t = s, A_t = a] = \max_{\pi} \mathbb{E}[R_{t+1} + \gamma G_{t+1} | S_t = s, A_t = a]$$
$$= \sum_{s', r} P(s', r | s, a) [r + \gamma \max_{a'} Q^*(s', a') | S_t = s, A_t = a] \text{ for all } s, a$$

Markov Decision Processes


Formal definition

Finite MDPs: (S, A, T, R)

1. $S_t \in S$ for $t = 0, 1, 2, 3, \dots$
2. $A_t \in A(s)$
3. $R_{t+1} \in R$
4. $T(s' | s, a) = \sum_{r' \in R} P(s', r' | s, a)$

Goal: maximize discounted **returns**

$$G_t = R_{t+1} + \gamma R_{t+2} + \gamma^2 R_{t+3} + \dots + \gamma^{T-t-1} R_T = \sum_{k=t+1}^T \gamma^{k-t-1} R_k$$
$$= R_{t+1} + \gamma G_{t+1}$$



- ▶ Optimization problem: (computationally) find **optimal policy** $\pi^*(S_t) = P(A_t | S_t)$
- ▶ Optimal policy $\pi^* : \pi \geq \pi' \Leftrightarrow v_{\pi^*}^*(s) \geq v_{\pi'}(s)$ for all s and $q_{\pi^*}^*(s, a) = \max_{\pi'} q_{\pi'}(s, a)$
- ▶ **Can be estimated from experience!**
 - e.g., value iteration methods for episode tasks

Multi-Armed Bandits

Finding the best restaurant

Finite MDPs: (S, A, T, R)


1. $S_t \in S$ where $S = \{s\}$

t is iteration over repeated choice now

2. $A(s) = \{a_0, a_1, \dots\}$

3. $R_t \in R$

4. $T(s' | s, a) = \sum_{r' \in R} P(s', r' | s, a)$



Goal: maximize reward over repeated action selection

Optimal **action-value function**:

$$Q_{\pi}^*(s, a) = \max_{\pi} \mathbb{E}[G_t | S_t = s, A_t = a]$$

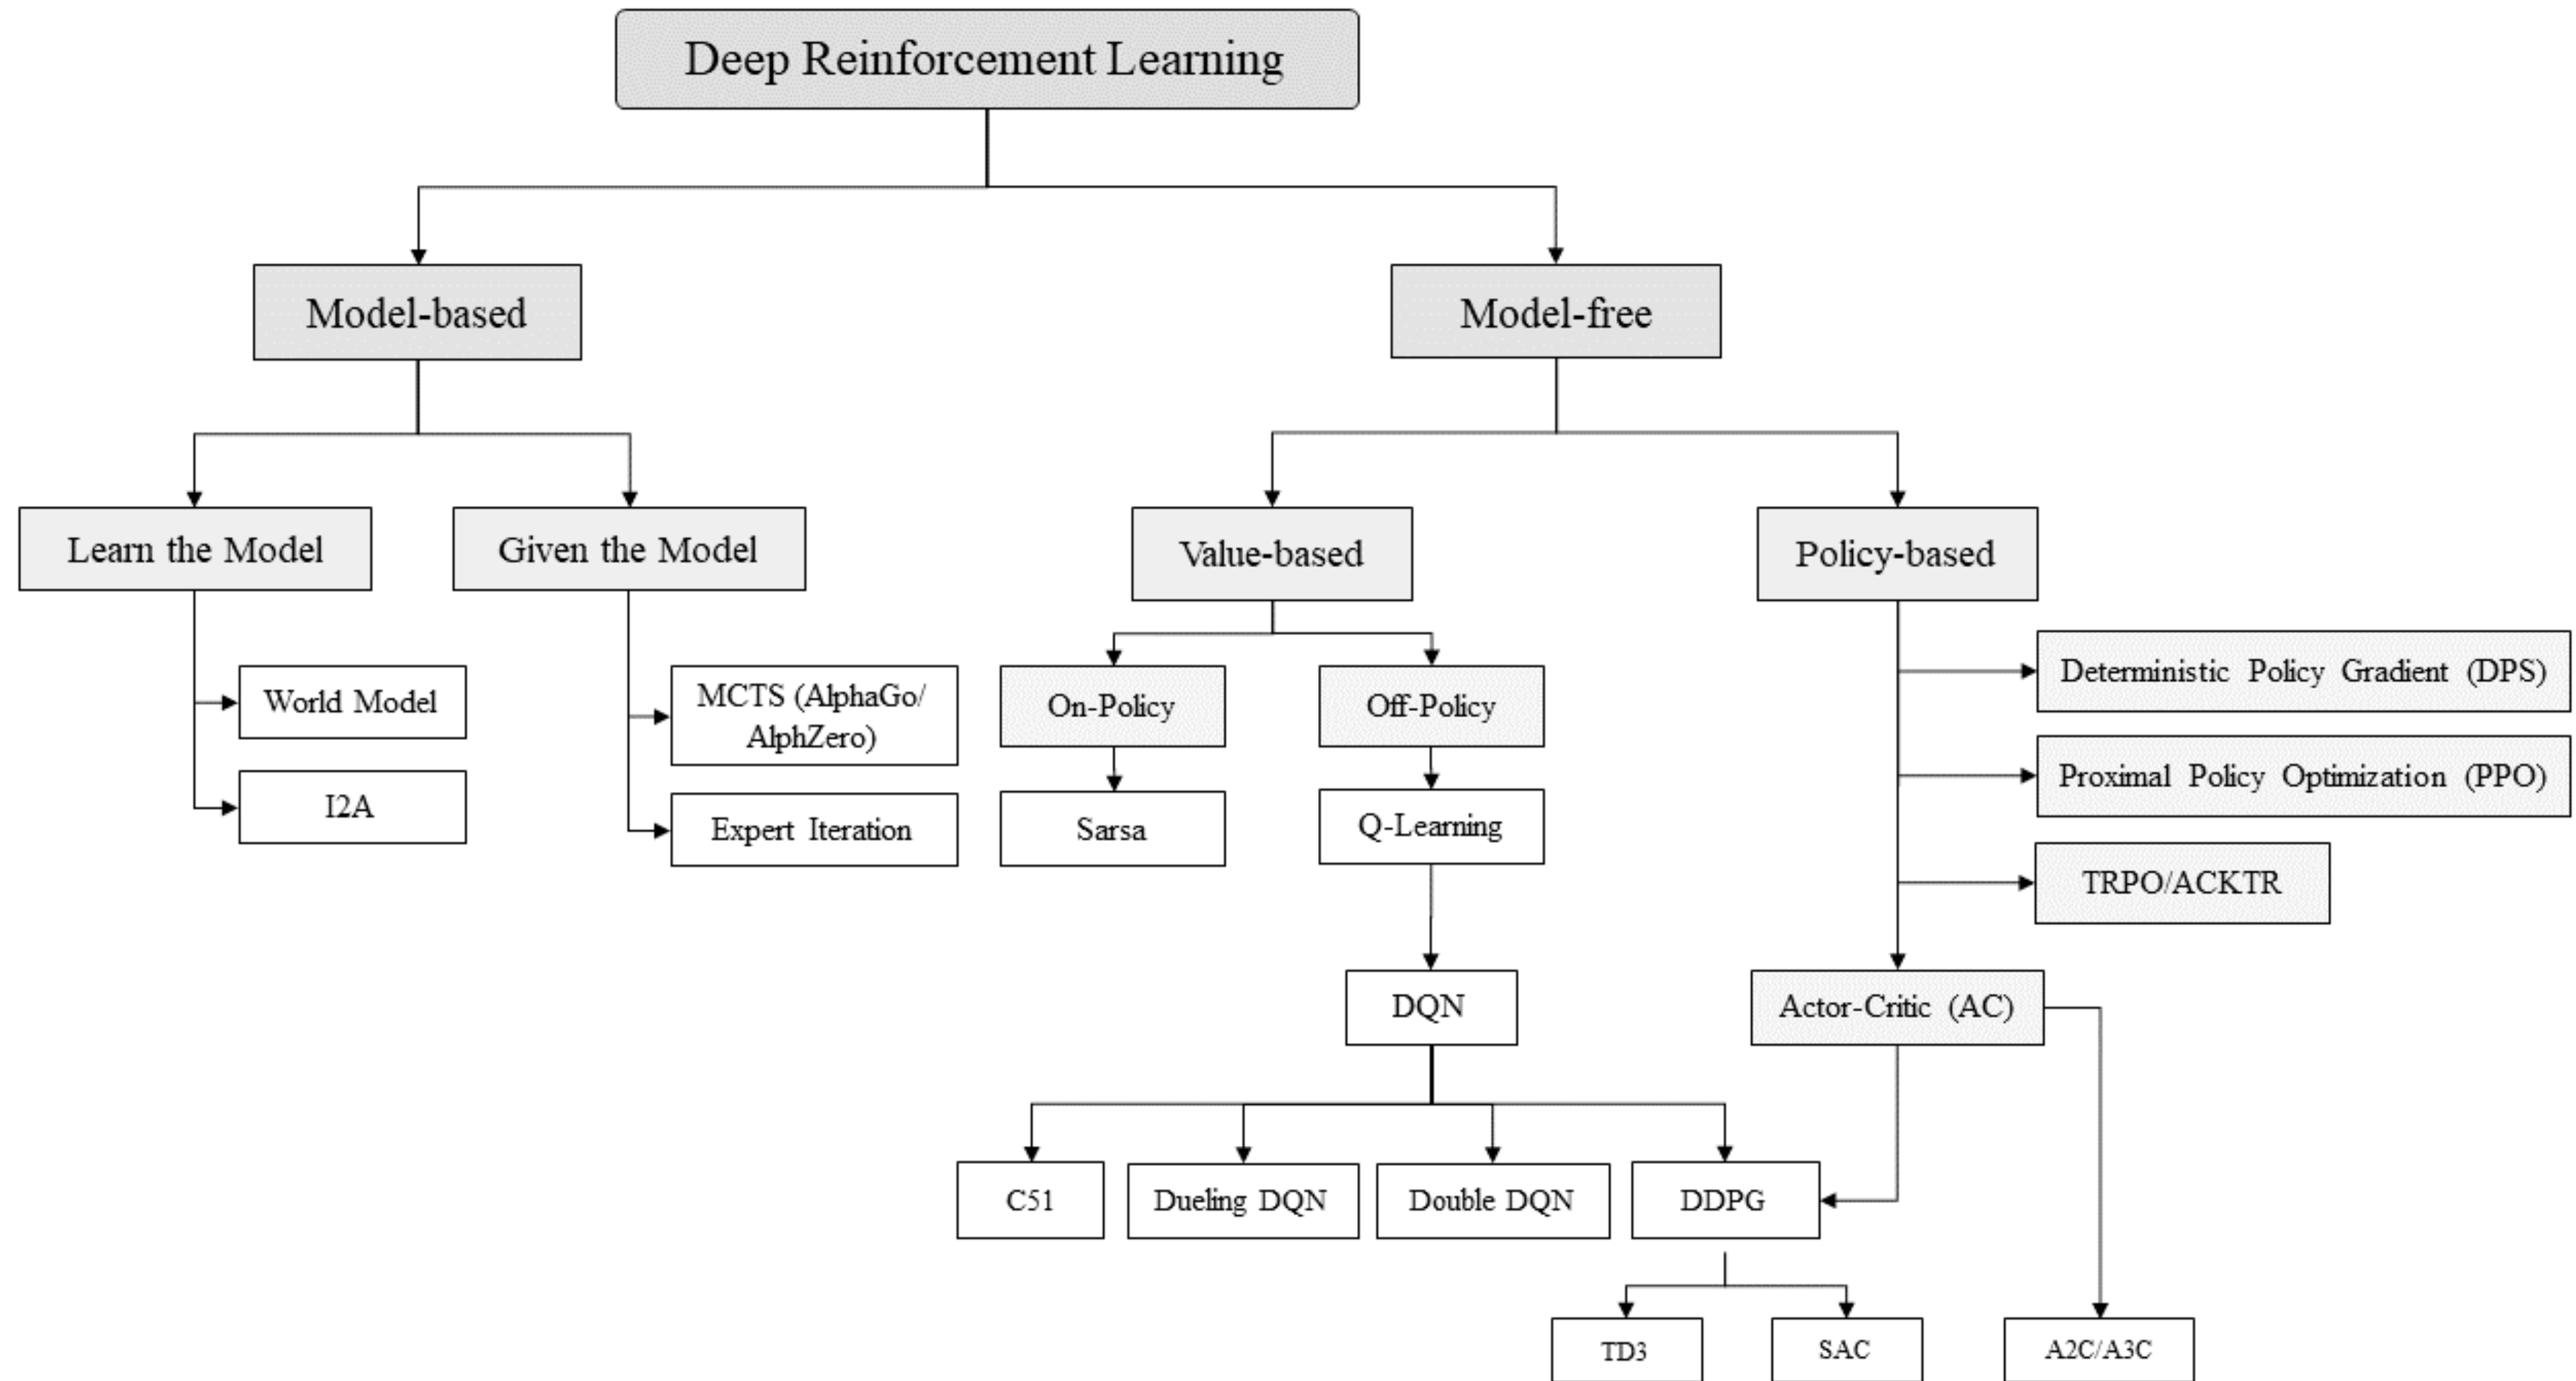


$$Q_{\pi}^*(a) = \max_{\pi} \mathbb{E}[R_t | A_t = a]$$

- Optimization problem: (computationally) find optimal policy $\pi^*(S_t) = P(A_t | S_t)$
- **Can be estimated from experience over repeated decision-making!**
 - e.g., with action-value methods like the sample-average method
- central problem: **exploration vs. exploitation**

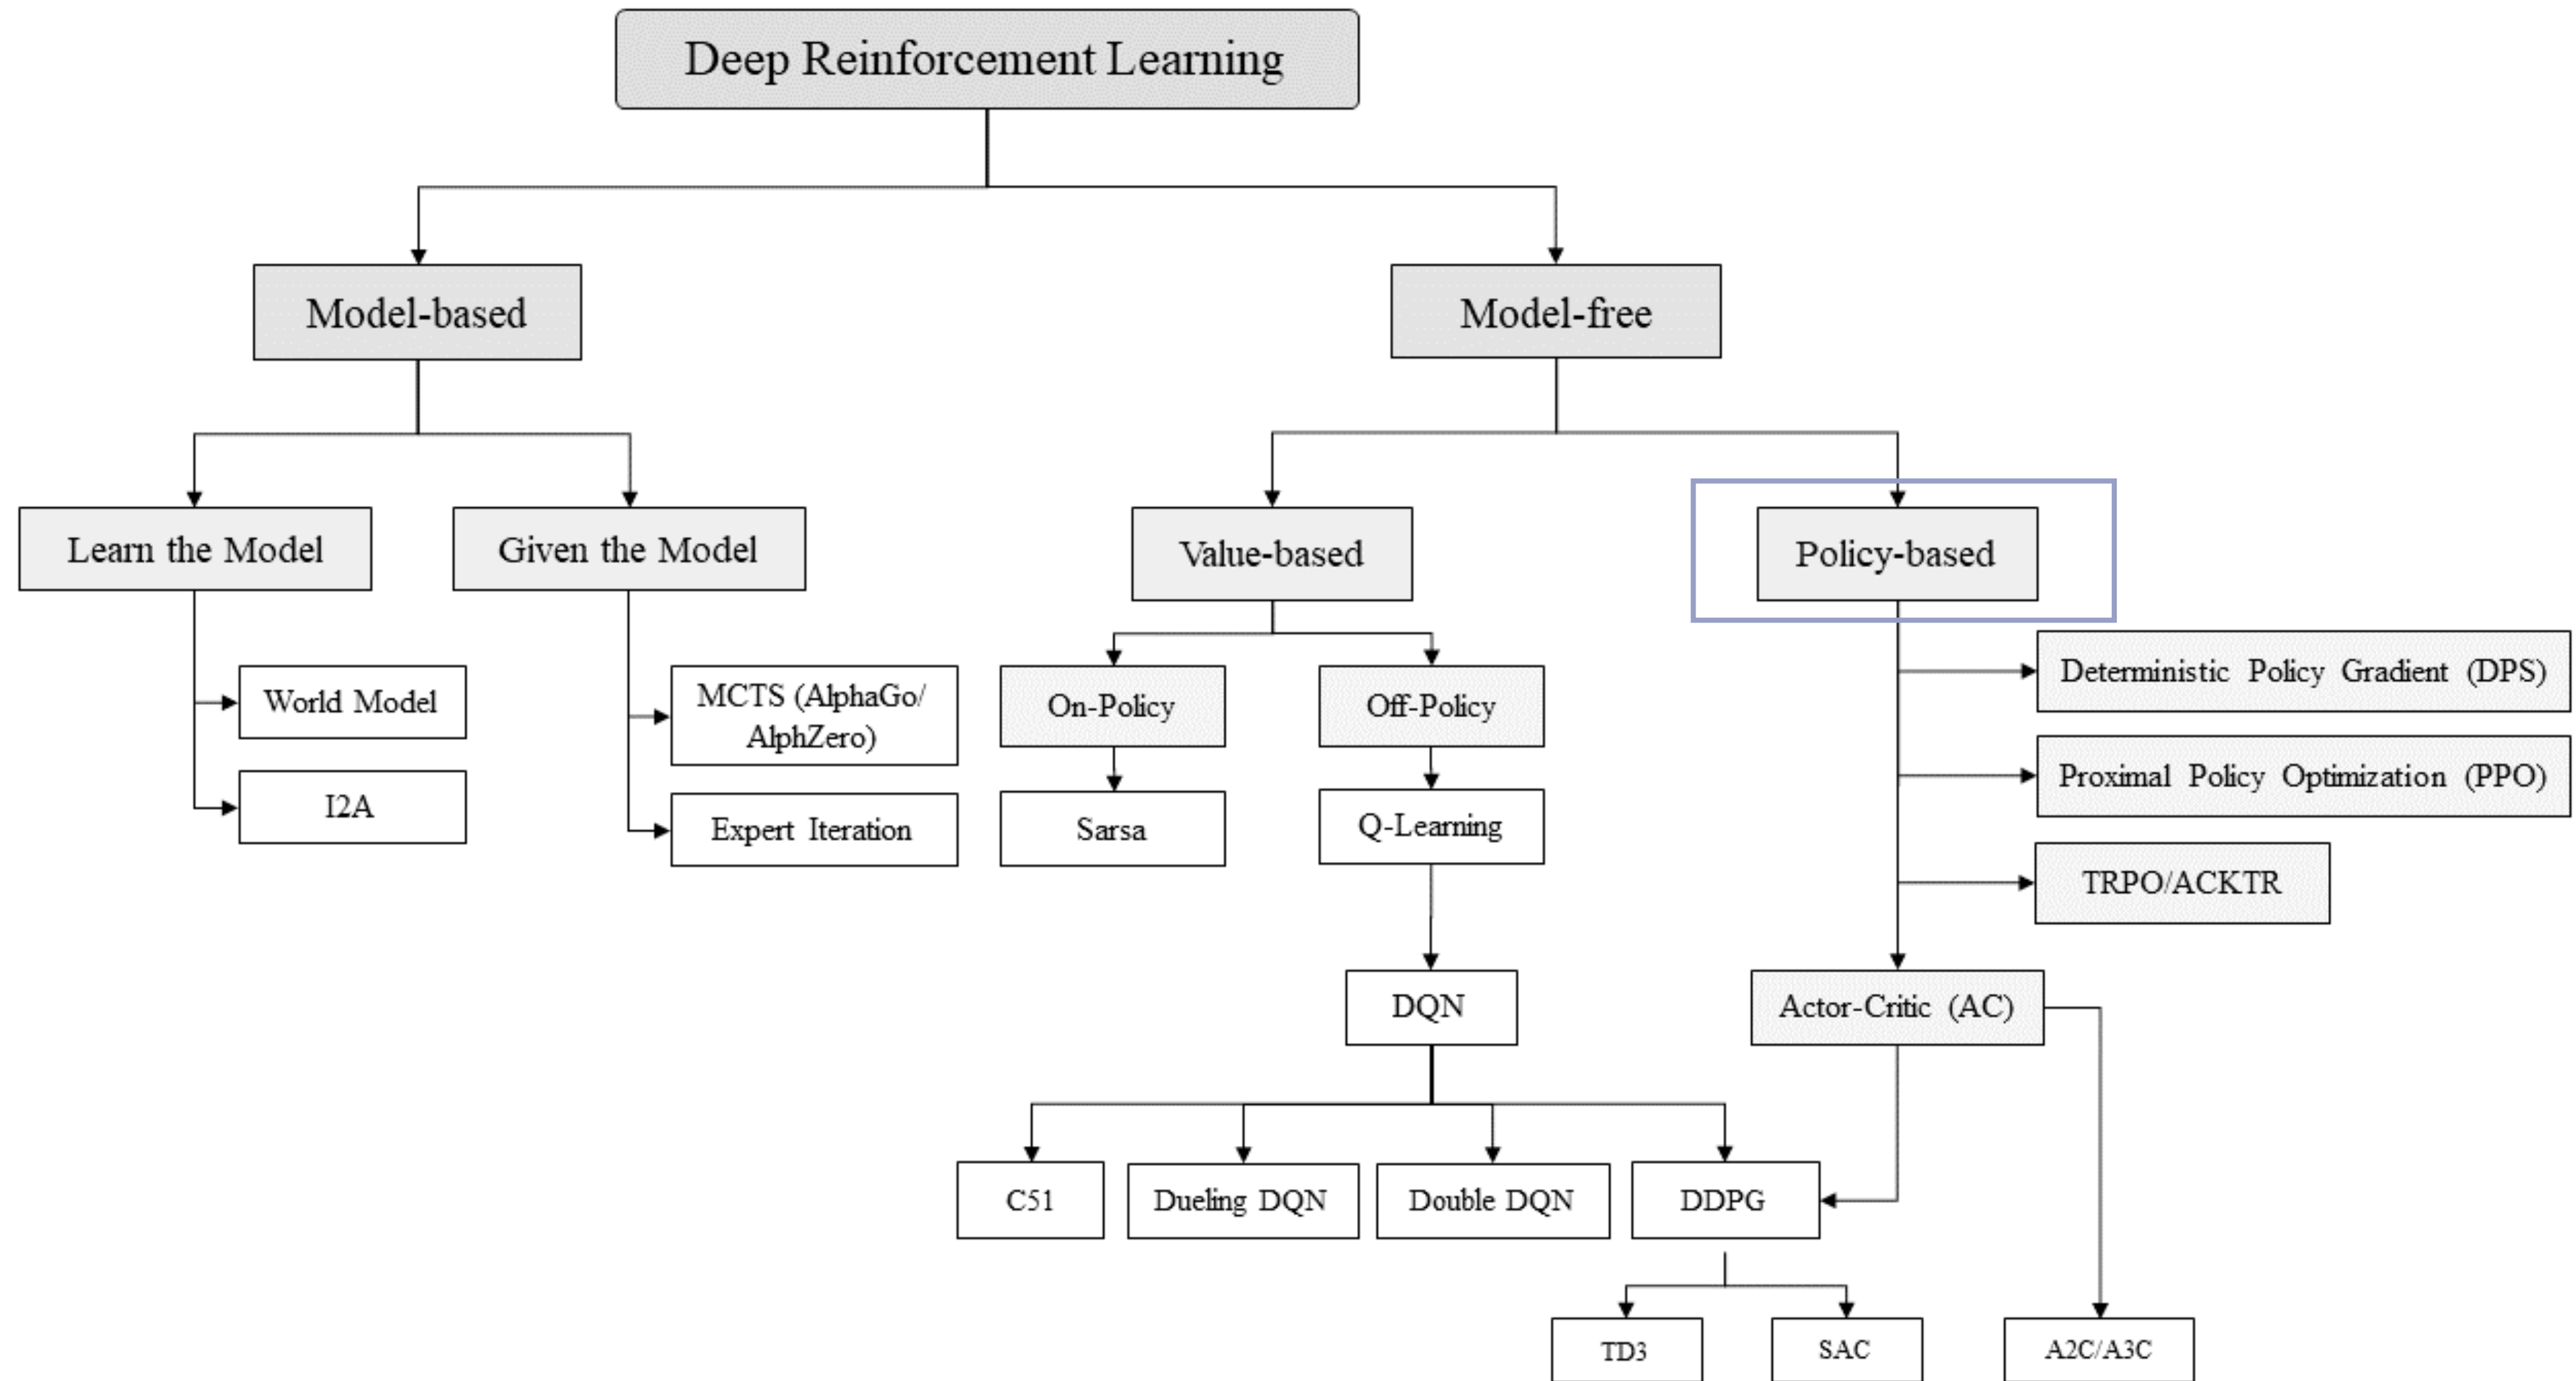
RL Algorithms

Approximating Optimal Policy



RL Algorithms

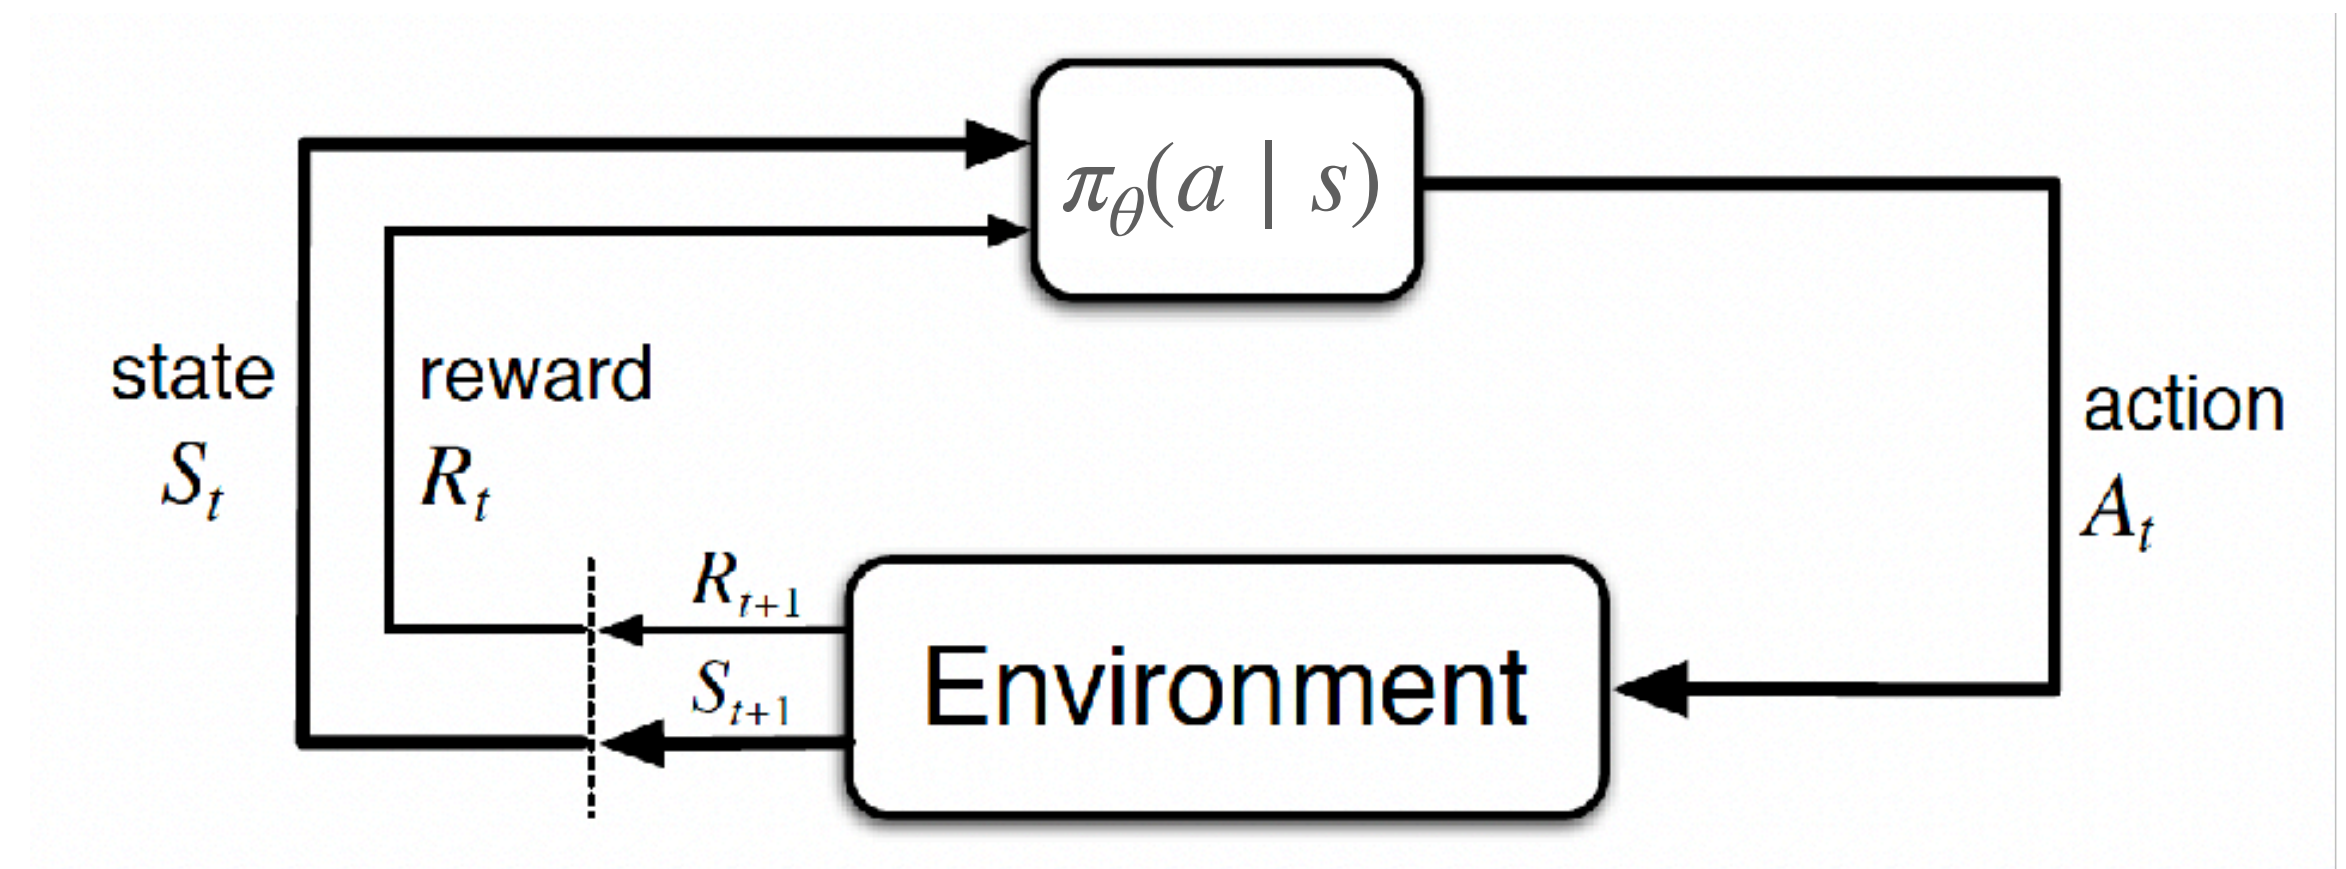
Approximating Optimal Policy



Policy-Gradient Methods

Introduction

- ▶ so far: deriving optimal policy from estimated value function
 - coming up with value functions might be difficult
 - state-value function doesn't prescribe actions
 - action-value functions require argmax
- ▶ idea: **optimize policy directly, such that expected reward is maximized**
 - think: optimize model with respect to objective function
- ▶ goal: find optimal θ
 - $\max_{\theta} \mathbb{E}_{\pi_{\theta}}[G_t]$
- ▶ recall LM optimization: tweak θ so as to minimize loss
 - Gradient descent: $\theta_{new} = \theta_{old} - \alpha \nabla L_{\theta}$
 - Now: gradient ascent: $\theta_{new} = \theta_{old} + \alpha \nabla L_{\theta}$



Policy-Gradient Methods

Policy-gradient theorem

- ▶ goal: find optimal θ
 - Now: gradient ascent: $\theta_{new} = \theta_{old} + \alpha \nabla L_{\theta}$
- ▶ we write τ for a sequence of states, actions, rewards and $R(\tau)$ for (discounted) return

$$L(\theta) = \sum_{\tau} P(\tau, \theta) R(\tau)$$

- ▶ sample-based policy gradient estimation

$$\begin{aligned} \nabla L(\theta) &= \nabla \sum_{\tau} P(\tau, \theta) R(\tau) = \sum_{\tau} \nabla_{\theta} P(\tau, \theta) R(\tau) \\ &= \sum_{\tau} \frac{P(\tau, \theta)}{P(\tau, \theta)} \nabla_{\theta} P(\tau, \theta) R(\tau) \\ &= \sum_{\tau} P(\tau, \theta) \frac{\nabla_{\theta} P(\tau, \theta)}{P(\tau, \theta)} R(\tau) = \sum_{\tau} P(\tau, \theta) \nabla_{\theta} \log P(\tau, \theta) R(\tau) \\ &\approx \frac{1}{m} \sum_{i=1}^m \nabla_{\theta} \log P(\tau^i, \theta) R(\tau^i) \end{aligned}$$

$$\nabla \log(f(x)) = \nabla f(x)/f(x)$$

Policy-Gradient Methods

REINFORCE

Policy gradient estimation: $\nabla L(\theta) = \sum_{\tau} P(\tau, \theta) \nabla_{\theta} \log P(\tau, \theta) R(\tau) \approx \frac{1}{m} \sum_{i=1}^m \nabla_{\theta} \log P(\tau^i, \theta) R(\tau^i)$

REINFORCE update rule:

for each episode:

generate $S_0, A_0, R_1, \dots, S_{T-1}, A_{T-1}, R_T$ following $\pi(\cdot | \cdot, \theta)$

for t in $\{0, 1, \dots, T-1\}$:

$$G \leftarrow \sum_{k=t+1}^T \gamma^{k-t-1} R_k$$

$$\theta_{t+1} \leftarrow \theta_t + \alpha G \nabla \log \pi(A_t | S_t, \theta_t)$$

Actor-Critic Algorithms

REINFORCE

(Advantage) Actor-Critic (A2C)

$$\theta_{t+1} = \theta_t + \alpha(G_t - \hat{v}(s, \mathbf{w})) \nabla \log \pi(A | s, \theta_t)$$

- ▶ Actor-critic methods evaluate the value of the action taken in a state:

$$\theta_{t+1} = \theta_t + \alpha(\boxed{G_{t:t+1}} - \boxed{\hat{v}(s, \mathbf{w})}) \nabla \log \boxed{\pi(A | s, \theta_t)}$$

One-step return Critic

Actor

▶

Policy-Gradient Methods

TRPO & PPO

$$\nabla L(\theta_t) \propto \mathbb{E}_{\pi}[G_t \nabla \log \pi(A | s, \theta)]$$

- Rewriting policy-gradient theorem:

$$L_{\theta} = \mathbb{E}_t[\hat{A}_t \log \pi_{\theta}(a_t | s_t)], \text{ where } \textbf{advantage } \hat{A}_t = G_t - b(S_t)$$

- Improve PG with Trust Region Policy Optimization (TRPO):

Surrogate objective

$$L^{CPI}(\theta) = \mathbb{E}_t\left[\frac{\pi_{\theta}(a_t | s_t)}{\pi_{\theta_old}(a_t | s_t)} \hat{A}_t\right] \text{ such that } \hat{\mathbb{E}}_t[KL[\pi_{\theta_old}(\cdot | s_t) \pi_{\theta}(\cdot | s_t)]] \leq \delta$$

Justification actually suggests penalty instead of constraint: $\mathbb{E}_t[r_t(\theta)A_t - \beta KL[\pi_{\theta_old}(\cdot | s_t) \pi_{\theta}(\cdot | s_t)]]$

Policy-Gradient Methods

TRPO & PPO

$$\nabla J(\theta_t) \propto \mathbb{E}_{\pi}[G_t \nabla \log \pi(A | s, \theta)]$$

- Rewriting policy-gradient theorem:

$$L_{\theta} = \mathbb{E}_t[\hat{A}_t \log \pi_{\theta}(a_t | s_t)], \text{ where advantage } \hat{A}_t = G_t - b(S_t)$$

- Improve PG with Trust Region Policy Optimization (TRPO):

Surrogate objective

$$L^{CPI}(\theta) = \mathbb{E}_t\left[\frac{\pi_{\theta}(a_t | s_t)}{\pi_{\theta_old}(a_t | s_t)} \hat{A}_t\right] \text{ such that } \mathbb{E}_t[KL[\pi_{\theta_old}(\cdot | s_t) \pi_{\theta}(\cdot | s_t)]] \leq \delta$$

Justification actually suggests penalty instead of constraint: $\mathbb{E}_t[r_t(\theta) \hat{A}_t - \beta KL[\pi_{\theta_old}(\cdot | s_t) \pi_{\theta}(\cdot | s_t)]]$

- Clip updates with PPO:

Expectation over batch

Boundaries

$$L^{CLIP}(\theta) = \mathbb{E}_t[\min(r_t(\theta) \hat{A}_t, \text{clip}(r_t(\theta), 1 - \epsilon, 1 + \epsilon) \hat{A}_t)]$$

Lower bound on
unclipped objective

Standard probability
ratio

Policy-Gradient Methods

Language models as policies

Policy gradient estimation: $\nabla L(\theta) = \sum_{\tau} P(\tau, \theta) \nabla_{\theta} \log P(\tau, \theta) R(\tau) \approx \frac{1}{m} \sum_{i=1}^m \nabla_{\theta} \log P(\tau^i, \theta) R(\tau^i)$

- ▶ policy P : language model
- ▶ trajectories τ : generations from language model
- ▶ $\log P(\tau^i, \theta)$: log probability of a generation τ^i under the language model
- ▶ $R(\tau^i)$: **reward for generation τ^i**

Summary

Reinforcement learning

- ▶ the central framework for formalizing RL problems are Markov Decision Processes (MDPs)
- ▶ task of RL is to solve MDP such that the expected return is maximized
 - and to find the optimal policy
- ▶ classical solution methods for MDPs include estimation of optimal state- and action-value functions
- ▶ policy gradient methods directly optimize the policy such that the expected return is maximized
 - can be applied to LMs!



Announcements

Solutions to exercises will be on Moodle on November 15th!

Next class (November 15th) online only!